

```
# IMPORTANT: RUN THIS CELL IN ORDER TO IMPORT YOUR KAGGLE DATA SOURCES,  
# THEN FEEL FREE TO DELETE THIS CELL.  
# NOTE: THIS NOTEBOOK ENVIRONMENT DIFFERS FROM KAGGLE'S PYTHON  
# ENVIRONMENT SO THERE MAY BE MISSING LIBRARIES USED BY YOUR  
# NOTEBOOK.  
import kagglehub  
urvishahir_electric_vehicle_specifications_dataset_2025_path = kagglehub.dataset_download('urvishahir/electric-vehicle-specifications-dataset-2025')  
  
print('Data source import complete.')
```

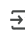
```
# This Python 3 environment comes with many helpful analytics libraries installed  
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python  
# For example, here's several helpful packages to load
```

```
import numpy as np # linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

```
# Input data files are available in the read-only "../input/" directory  
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory
```

```
import os  
for dirname, _, filenames in os.walk('/kaggle/input'):  
    for filename in filenames:  
        print(os.path.join(dirname, filename))
```

```
# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run A  
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

```
 /kaggle/input/electric-vehicle-specifications-dataset-2025/electric_vehicles_spec_2025.csv.csv
```

Double-click (or enter) to edit

```
# Import libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from scipy import stats  
import plotly.express as px  
import plotly.graph_objects as go  
from plotly.subplots import make_subplots
```

```
df = pd.read_csv("/kaggle/input/electric-vehicle-specifications-dataset-2025/electric_vehicles_spec_2025.csv.csv")  
df
```

```

/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1458: RuntimeWarning: invalid value encountered in greater
has_large_values = (abs_vals > 1e6).any()
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in less
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in greater
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1458: RuntimeWarning: invalid value encountered in greater
has_large_values = (abs_vals > 1e6).any()
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in less
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in greater
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()

```

	brand	model	top_speed_kmh	battery_capacity_kWh	battery_type	number_of_cells	torque_nm	efficiency_wh_per_km	range_km	acceler
0	Abarth	500e Convertible	155	37.8	Lithium-ion	192.0	235.0	156	225	
1	Abarth	500e Hatchback	155	37.8	Lithium-ion	192.0	235.0	149	225	
2	Abarth	600e Scorpionissima	200	50.8	Lithium-ion	102.0	345.0	158	280	
3	Abarth	600e Turismo	200	50.8	Lithium-ion	102.0	345.0	158	280	
4	Aiways	U5	150	60.0	Lithium-ion	NaN	310.0	156	315	
...
473	Zeekr	7X Premium RWD	210	71.0	Lithium-ion	NaN	440.0	148	365	
474	Zeekr	X Core RWD (MY25)	190	49.0	Lithium-ion	NaN	343.0	148	265	
475	Zeekr	X Long Range RWD (MY25)	190	65.0	Lithium-ion	NaN	343.0	146	360	
476	Zeekr	X Privilege AWD (MY25)	190	65.0	Lithium-ion	NaN	543.0	153	350	
477	firefly	NaN	150	41.2	Lithium-ion	112.0	200.0	125	250	

478 rows x 22 columns

▼ Data exploration

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 478 entries, 0 to 477
Data columns (total 22 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   brand                                478 non-null    object
 1   model                                477 non-null    object
 2   top_speed_kmh                        478 non-null    int64
 3   battery_capacity_kWh                 478 non-null    float64
 4   battery_type                         478 non-null    object
 5   number_of_cells                     276 non-null    float64
 6   torque_nm                           471 non-null    float64
 7   efficiency_wh_per_km                478 non-null    int64
 8   range_km                            478 non-null    int64
 9   acceleration_0_100_s                478 non-null    float64
10   fast_charging_power_kw_dc           477 non-null    float64
11   fast_charge_port                     477 non-null    object
12   towing_capacity_kg                  452 non-null    float64
13   cargo_volume_l                      477 non-null    object
14   seats                                478 non-null    int64
15   drivetrain                           478 non-null    object
16   segment                              478 non-null    object
17   length_mm                           478 non-null    int64
18   width_mm                            478 non-null    int64
19   height_mm                           478 non-null    int64
20   car_body_type                       478 non-null    object
21   source_url                          478 non-null    object
dtypes: float64(6), int64(7), object(9)
memory usage: 82.3+ KB

```

```
df.describe()
```

	top_speed_kmh	battery_capacity_kWh	number_of_cells	torque_nm	efficiency_wh_per_km	range_km	acceleration_0_100_s	fast_charging_pow
count	478.000000	478.000000	276.000000	471.000000	478.000000	478.000000	478.000000	4
mean	185.487448	74.043724	485.293478	498.012739	162.903766	393.179916	6.882636	1
std	34.252773	20.331058	1210.819733	241.461128	34.317532	103.287335	2.730696	
min	125.000000	21.300000	72.000000	113.000000	109.000000	135.000000	2.200000	
25%	160.000000	60.000000	150.000000	305.000000	143.000000	320.000000	4.800000	
50%	180.000000	76.150000	216.000000	430.000000	155.000000	397.500000	6.600000	1
75%	201.000000	90.600000	324.000000	679.000000	177.750000	470.000000	8.200000	1
max	325.000000	118.000000	7920.000000	1350.000000	370.000000	685.000000	19.100000	2

```
df.isnull().sum()
```

brand	0
model	1
top_speed_kmh	0
battery_capacity_kWh	0
battery_type	0
number_of_cells	202
torque_nm	7
efficiency_wh_per_km	0
range_km	0
acceleration_0_100_s	0
fast_charging_power_kw_dc	1
fast_charge_port	1
towing_capacity_kg	26
cargo_volume_l	1
seats	0
drivetrain	0
segment	0
length_mm	0
width_mm	0
height_mm	0
car_body_type	0
source_url	0
dtype: int64	


✦ Lis of brands

```
# Count brands
df['brand'].unique()

array(['Abarth', 'Aiways', 'Alfa', 'Alpine', 'Audi', 'BMW', 'BYD',
      'CUPRA', 'Cadillac', 'Citroen', 'DS', 'Dacia', 'Dongfeng',
      'Elaris', 'Fiat', 'Ford', 'GWM', 'Genesis', 'Honda', 'Hongqi',
      'Hyundai', 'Jaguar', 'Jeep', 'KGM', 'Kia', 'Lancia', 'Leapmotor',
      'Lexus', 'Lotus', 'Lucid', 'Lynk&Co', 'MG', 'Maserati', 'Maxus',
      'Mazda', 'Mercedes-Benz', 'Mini', 'NIO', 'Nissan', 'Omoda', 'Opel',
      'Peugeot', 'Polestar', 'Porsche', 'Renault', 'Rolls-Royce',
      'Skoda', 'Skywell', 'Smart', 'Subaru', 'Tesla', 'Toyota',
      'VinFast', 'Volkswagen', 'Volvo', 'Voyah', 'XPENG', 'Zeekr',
      'firefly'], dtype=object)
```

```
numerical_cols = df.select_dtypes(include='number').columns.tolist()
print("numerical_cols: ", numerical_cols)
```

```
sns.pairplot(df,
              vars=numerical_cols)
plt.suptitle("Pairplot of Selected Features", y=1.02)
plt.show()
```

✦ Correlation


```
correlation_matrix = df[numerical_cols].corr()

plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, fmt='.2f', cbar_kws={'label': 'Correlation Coefficient'})
plt.title('Correlation Matrix - Numerical Variables')
plt.tight_layout()
plt.show()

# Find highly correlated pairs
print("\nHighly correlated variable pairs (|correlation| > 0.7):")
high_corr_pairs = []
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        corr_val = correlation_matrix.iloc[i, j]
        if abs(corr_val) > 0.7:
            high_corr_pairs.append((
                correlation_matrix.columns[i],
                correlation_matrix.columns[j],
                corr_val
            ))

if high_corr_pairs:
    for var1, var2, corr in high_corr_pairs:
        print(f" {var1} <-> {var2}: {corr:.3f}")
```

```
numerical_cols: ['top_speed_kmh', 'battery_capacity_kWh', 'number_of_cells', 'torque_nm', 'efficiency_wh_per_kmh', 'range_km', 'acceleration_0_100_s']
/usr/local/lib/python3.11/dist-packages/matplotlib/colormaps.py:721: RuntimeWarning: invalid value encountered in less
  xa[xa < 0] = -1
```



Highly correlated variable pairs ($|\text{correlation}| > 0.7$):

```
top_speed_kmh <=> battery_capacity_kWh: 0.708
top_speed_kmh <=> torque_nm: 0.806
top_speed_kmh <=> range_km: 0.732
top_speed_kmh <=> acceleration_0_100_s: -0.823
top_speed_kmh <=> fast_charging_power_kw_dc: 0.772
battery_capacity_kWh <=> torque_nm: 0.757
battery_capacity_kWh <=> range_km: 0.880
battery_capacity_kWh <=> fast_charging_power_kw_dc: 0.728
battery_capacity_kWh <=> width_mm: 0.731
torque_nm <=> acceleration_0_100_s: -0.788
range_km <=> acceleration_0_100_s: -0.712
range_km <=> fast_charging_power_kw_dc: 0.721
length_mm <=> width_mm: 0.850
```

✓ Data visualization

✓ Distribution of Electric Vehicles by Brand

```
brand_counts = df['brand'].value_counts().head(15)
print(f"Top 15 brands by number of models:")
print(brand_counts)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
```

```
# Bar plot
brand_counts.plot(kind='bar', ax=ax1, color='skyblue')
ax1.set_title('Top 15 EV Brands by Number of Models', fontsize=14, fontweight='bold')
ax1.set_xlabel('Brand')
ax1.set_ylabel('Number of Models')
ax1.tick_params(axis='x', rotation=45)

# Pie chart for top 10
brand_counts.head(10).plot(kind='pie', ax=ax2, autopct='%1.1f%%')
ax2.set_title('Top 10 Brands - Market Share', fontsize=14, fontweight='bold')
ax2.set_ylabel('')


plt.tight_layout()
plt.show()
```

↗ Top 15 brands by number of models:


brand	
Mercedes-Benz	42
Audi	28
Porsche	26
Volkswagen	23
Ford	22
BMW	20
Peugeot	19
Volvo	18
BYD	17
Smart	17
Kia	16
Opel	14
Hyundai	14
Skoda	12
MG	12

Name: count, dtype: int64

Top 15 EV Brands by Number of Models



Top 10 Brands - Market Share



✓ Battery Capacity vs Driving Range Correlation

```
# Remove outliers for better visualization
df_clean = df.dropna(subset=['battery_capacity_kWh', 'range_km'])
df_clean = df_clean[(df_clean['battery_capacity_kWh'] < 200) & (df_clean['range_km'] < 1000)]

correlation = df_clean['battery_capacity_kWh'].corr(df_clean['range_km'])
print(f"Correlation coefficient: {correlation:.3f}")


plt.figure(figsize=(12, 6))

plt.scatter(df_clean['battery_capacity_kWh'], df_clean['range_km'], alpha=0.6, color='green')
plt.xlabel('Battery Capacity (kWh)')
plt.ylabel('Range (km)')
plt.title('Battery Capacity vs Range')
z = np.polyfit(df_clean['battery_capacity_kWh'], df_clean['range_km'], 1)
p = np.poly1d(z)
plt.plot(df_clean['battery_capacity_kWh'], p(df_clean['battery_capacity_kWh']), "r--", alpha=0.8)

plt.tight_layout()
```

```
plt.show()
```

Correlation coefficient: 0.880



What are the most common drivetrain types?

```
drivetrain_counts = df['drivetrain'].value_counts()
print("Drivetrain distribution:")
print(drivetrain_counts)

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
drivetrain_counts.plot(kind='bar', color=['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4'])
plt.title('Drivetrain Types Distribution')
plt.xlabel('Drivetrain Type')
plt.ylabel('Number of Vehicles')
plt.xticks(rotation=45)

plt.subplot(1, 2, 2)
plt.pie(drivetrain_counts.values, labels=drivetrain_counts.index, autopct='%1.1f%%', startangle=90)
plt.title('Drivetrain Types - Percentage Distribution')


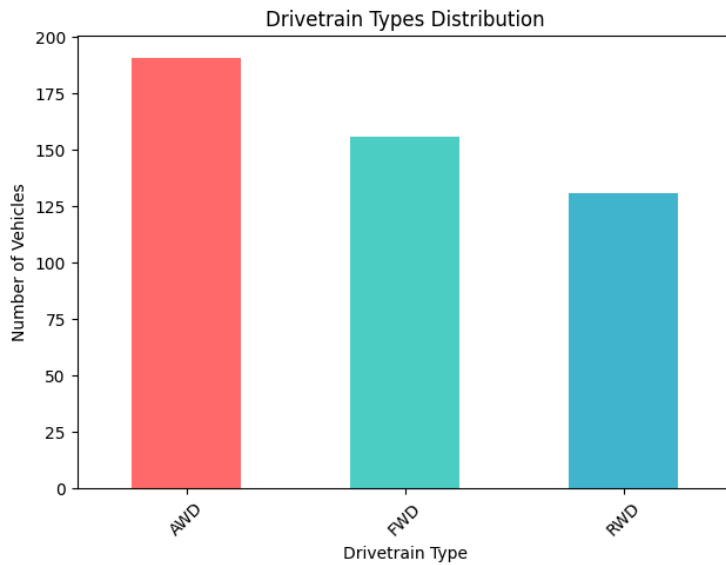
plt.tight_layout()
plt.show()
```



```

Drivetrain distribution:
drivetrain
AWD    191
FWD    156
RWD    131
Name: count, dtype: int64

```



How does acceleration performance vary by segment?

```

df_accel = df.dropna(subset=['acceleration_0_100_s', 'segment'])
df_accel = df_accel[df_accel['acceleration_0_100_s'] < 20] # Remove outliers

print("Average acceleration by segment:")
segment_accel = df_accel.groupby('segment')['acceleration_0_100_s'].agg(['mean', 'count']).round(2)
print(segment_accel)

plt.figure(figsize=(15, 6))

plt.subplot(1, 2, 1)
sns.boxplot(data=df_accel, x='segment', y='acceleration_0_100_s')
plt.title('Acceleration Performance by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('0-100 km/h (seconds)')
plt.xticks(rotation=45)


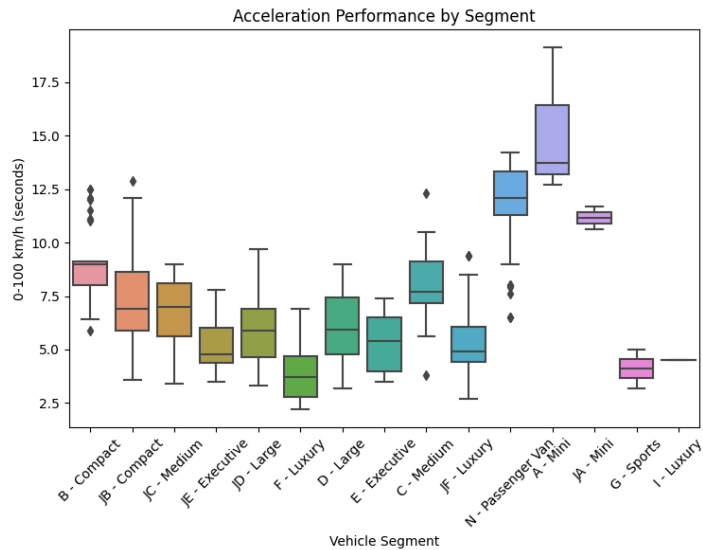
plt.subplot(1, 2, 2)
segment_means = df_accel.groupby('segment')['acceleration_0_100_s'].mean().sort_values()
segment_means.plot(kind='bar', color='orange')
plt.title('Average Acceleration by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Average 0-100 km/h (seconds)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

```

➡ Average acceleration by segment:

	mean	count
segment		
A - Mini	15.17	3
B - Compact	8.99	29
C - Medium	7.98	34
D - Large	6.00	28
E - Executive	5.25	30
F - Luxury	3.89	51
G - Sports	4.10	2
I - Luxury	4.50	1
JA - Mini	11.15	2
JB - Compact	7.16	44
JC - Medium	6.94	91
JD - Large	6.03	58
JE - Executive	5.13	28
JF - Luxury	5.41	30
N - Passenger Van	11.74	47



✓ What is the relationship between top speed and acceleration?

```
df_perf = df.dropna(subset=['top_speed_kmh', 'acceleration_0_100_s'])
df_perf = df_perf[(df_perf['top_speed_kmh'] < 300) & (df_perf['acceleration_0_100_s'] < 15)]

correlation = df_perf['top_speed_kmh'].corr(df_perf['acceleration_0_100_s'])
print(f"Correlation between top speed and acceleration: {correlation:.3f}")

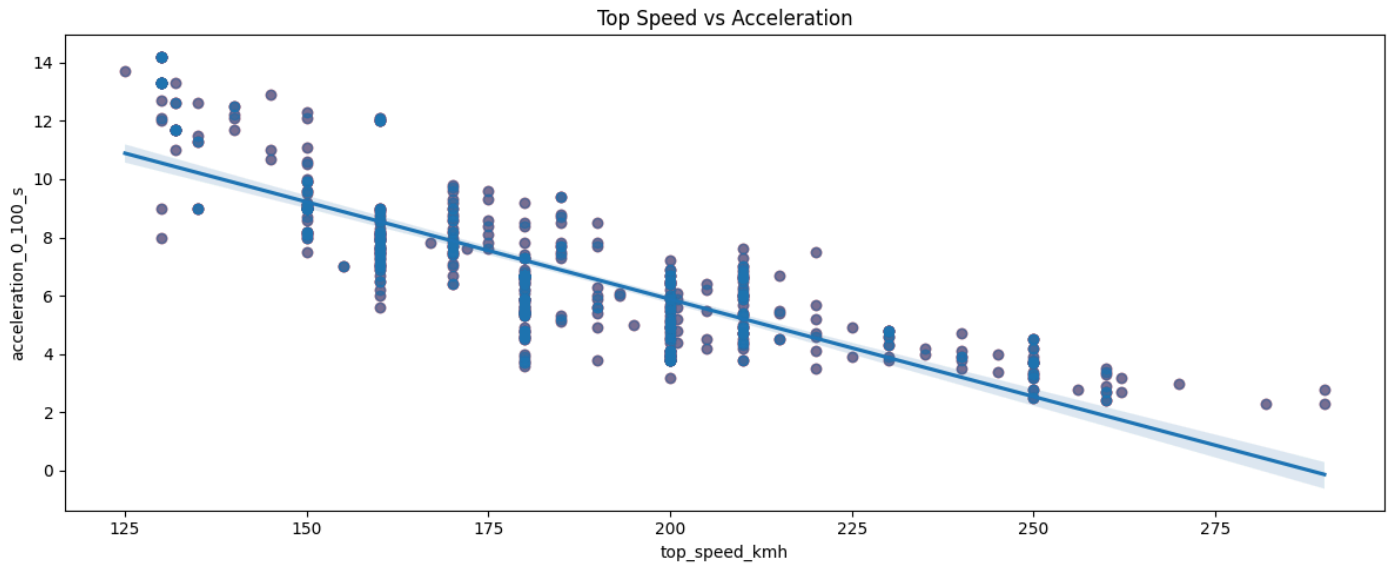
plt.figure(figsize=(12, 5))

plt.scatter(df_perf['top_speed_kmh'], df_perf['acceleration_0_100_s'], alpha=0.6, color='red')
plt.xlabel('Top Speed (km/h)')
plt.ylabel('0-100 km/h (seconds)')
plt.title('Top Speed vs Acceleration')

sns.regplot(data=df_perf, x='top_speed_kmh', y='acceleration_0_100_s', scatter_kws={'alpha':0.6})

plt.tight_layout()
plt.show()
```

↔ Correlation between top speed and acceleration: -0.831



Vehicle Distribution Across Segments

segment_counts

```
↔ segment
JC - Medium      91
JD - Large       58
F - Luxury       51
N - Passenger Van 47
JB - Compact     44
C - Medium       34
E - Executive    30
JF - Luxury      30
B - Compact      29
JE - Executive   28
D - Large        28
A - Mini         3
JA - Mini        2
G - Sports       2
I - Luxury       1
Name: count, dtype: int64
```

```
segment_counts = df['segment'].value_counts()
print("Segment distribution:")
print(segment_counts)
```

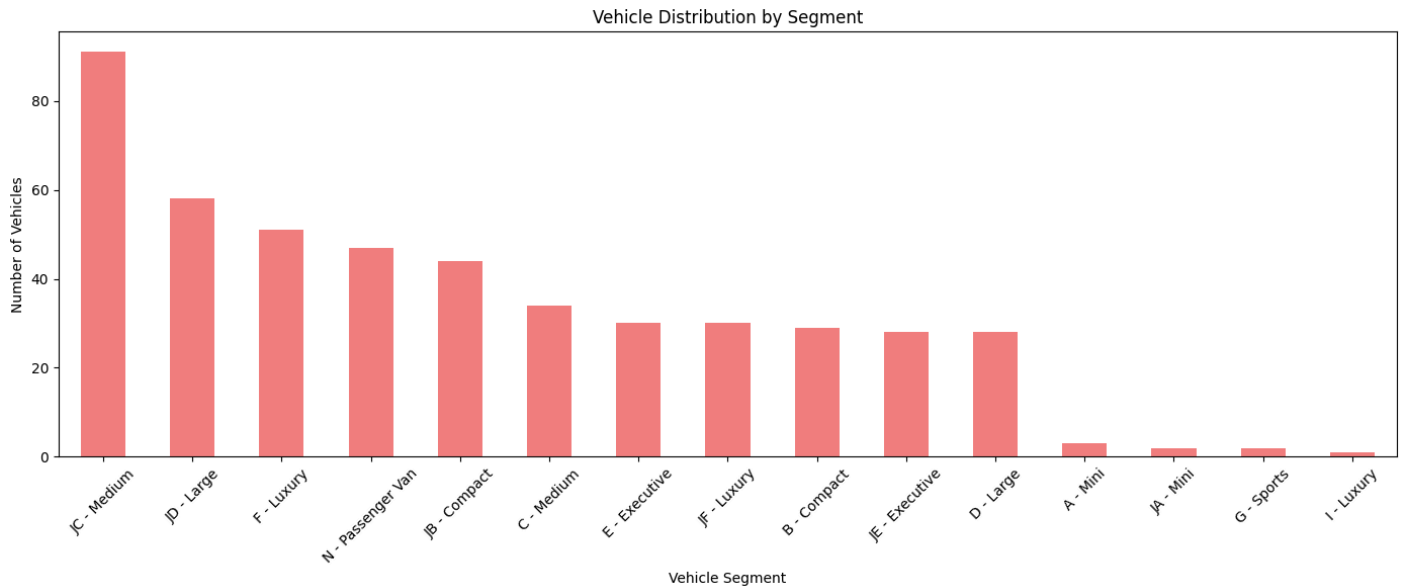
```
plt.figure(figsize=(14, 6))
```

```
segment_counts.plot(kind='bar', color='lightcoral')
plt.title('Vehicle Distribution by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Number of Vehicles')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

```

↔ Segment distribution:
segment
JC - Medium      91
JD - Large       58
F - Luxury       51
N - Passenger Van 47
JB - Compact     44
C - Medium       34
E - Executive    30
JF - Luxury      30
B - Compact      29
JE - Executive   28
D - Large        28
A - Mini         3
JA - Mini        2
G - Sports       2
I - Luxury       1
Name: count, dtype: int64

```



✓ What body types are most common in the EV market?

```

body_counts = df['car_body_type'].value_counts()
print("Body type distribution:")
print(body_counts)

```

```

plt.figure(figsize=(12, 6))

```

```

body_counts.plot(kind='bar', color='lightgreen')
plt.title('Body Types Distribution')
plt.xlabel('Body Type')
plt.ylabel('Number of Vehicles')
plt.xticks(rotation=45)

```

```

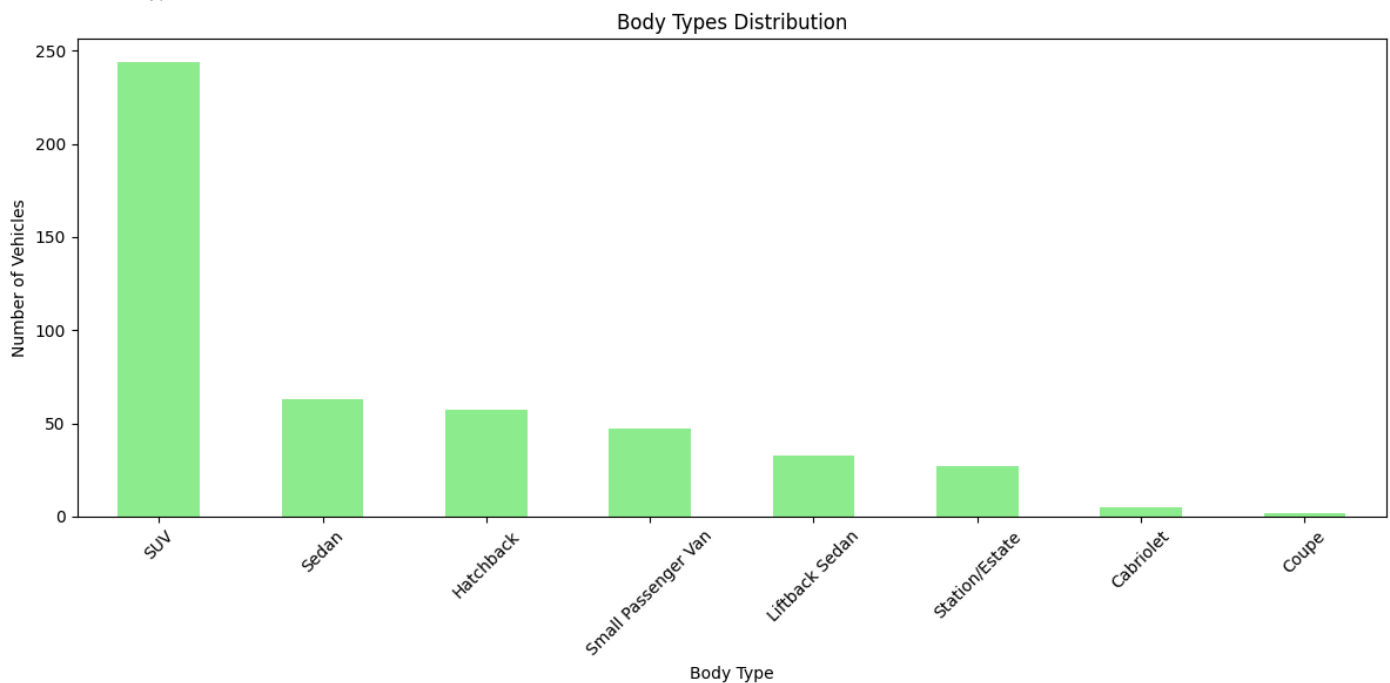
plt.tight_layout()
plt.show()

```

```

Body type distribution:
car_body_type
SUV                244
Sedan              63
Hatchback          57
Small Passenger Van 47
Liftback Sedan     33
Station/Estate     27
Cabriolet          5
Coupe              2
Name: count, dtype: int64

```



✓ How does battery capacity vary by vehicle segment?

```

df_battery = df.dropna(subset=['battery_capacity_kWh', 'segment'])
df_battery = df_battery[df_battery['battery_capacity_kWh'] < 200] # Remove outliers

print("Battery capacity statistics by segment:")
battery_stats = df_battery.groupby('segment')['battery_capacity_kWh'].agg(['mean', 'median', 'std', 'count']).round(2)
print(battery_stats)

plt.figure(figsize=(15, 6))

plt.subplot(1, 2, 1)
sns.boxplot(data=df_battery, x='segment', y='battery_capacity_kWh')
plt.title('Battery Capacity Distribution by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Battery Capacity (kWh)')
plt.xticks(rotation=45)

plt.subplot(1, 2, 2)
segment_battery_mean = df_battery.groupby('segment')['battery_capacity_kWh'].mean().sort_values(ascending=False)
segment_battery_mean.plot(kind='bar', color='purple')
plt.title('Average Battery Capacity by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Average Battery Capacity (kWh)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

```

```

Battery capacity statistics by segment:
      mean  median  std  count
segment

```

segment	mean	median	std	count
A - Mini	28.67	25.00	6.35	3
B - Compact	40.47	41.20	8.94	29
C - Medium	59.14	58.20	11.62	34
D - Large	73.43	75.00	9.60	28
E - Executive	85.92	86.00	6.41	30
F - Luxury	96.59	97.00	9.97	51
G - Sports	74.40	74.40	0.00	2
I - Luxury	102.00	102.00	NaN	1
JA - Mini	42.50	42.50	4.95	2
JB - Compact	54.43	50.80	8.84	44
JC - Medium	70.85	73.00	9.56	91
JD - Large	84.58	87.85	10.81	58
JE - Executive	94.23	94.90	10.01	28
JF - Luxury	100.65	104.00	13.85	30
N - Passenger Van	60.01	60.00	14.49	47

```

/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1458: RuntimeWarning: invalid value encountered in greater
has_large_values = (abs_vals > 1e6).any()

```

```

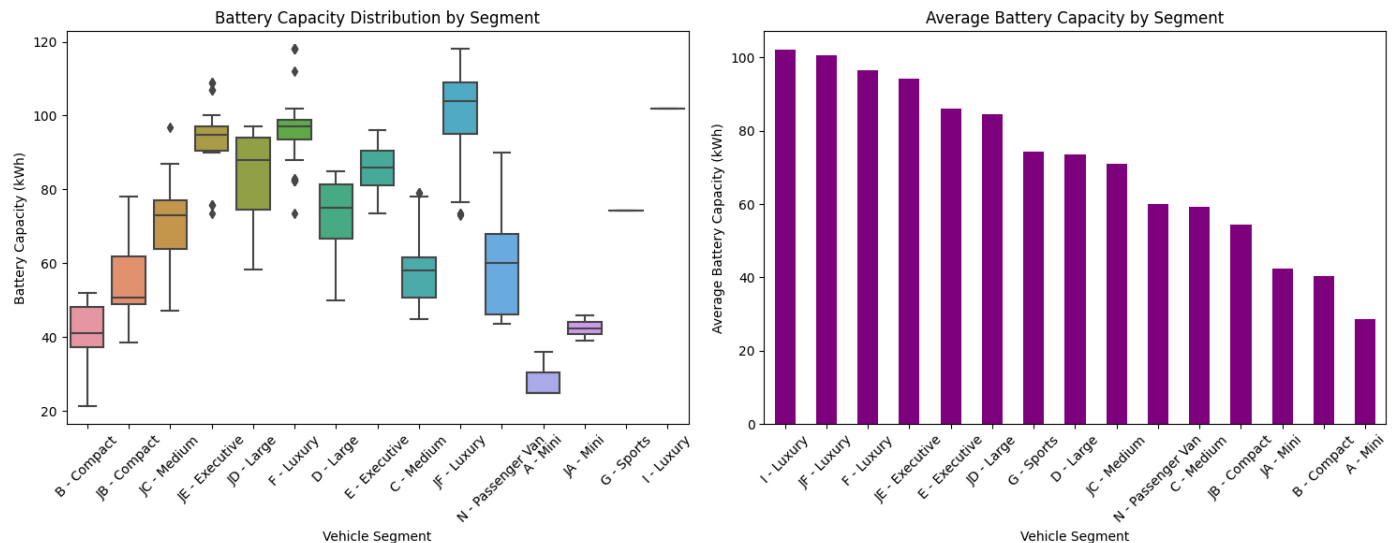
/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in less
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()

```

```

/usr/local/lib/python3.11/dist-packages/pandas/io/formats/format.py:1459: RuntimeWarning: invalid value encountered in greater
has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals > 0)).any()

```



What is the average range by drivetrain type?

```

df_range = df.dropna(subset=['range_km', 'drivetrain'])
df_range = df_range[df_range['range_km'] < 1000] # Remove outliers

print("Range statistics by drivetrain:")
range_stats = df_range.groupby('drivetrain')['range_km'].agg(['mean', 'median', 'std', 'count']).round(2)
print(range_stats)

plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
sns.boxplot(data=df_range, x='drivetrain', y='range_km')
plt.title('Range Distribution by Drivetrain')
plt.xlabel('Drivetrain Type')
plt.ylabel('Range (km)')

plt.subplot(1, 2, 2)
drivetrain_range_mean = df_range.groupby('drivetrain')['range_km'].mean().sort_values(ascending=False)
drivetrain_range_mean.plot(kind='bar', color='teal')
plt.title('Average Range by Drivetrain')
plt.xlabel('Drivetrain Type')
plt.ylabel('Average Range (km)')

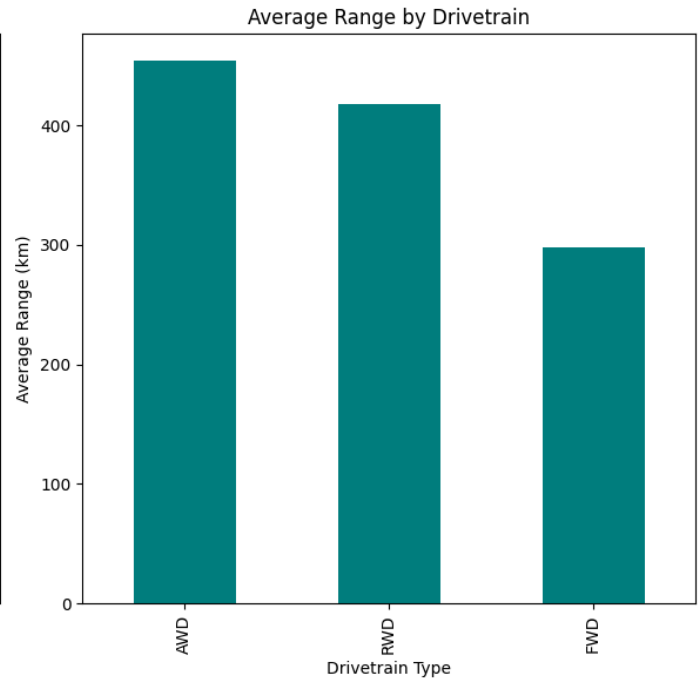
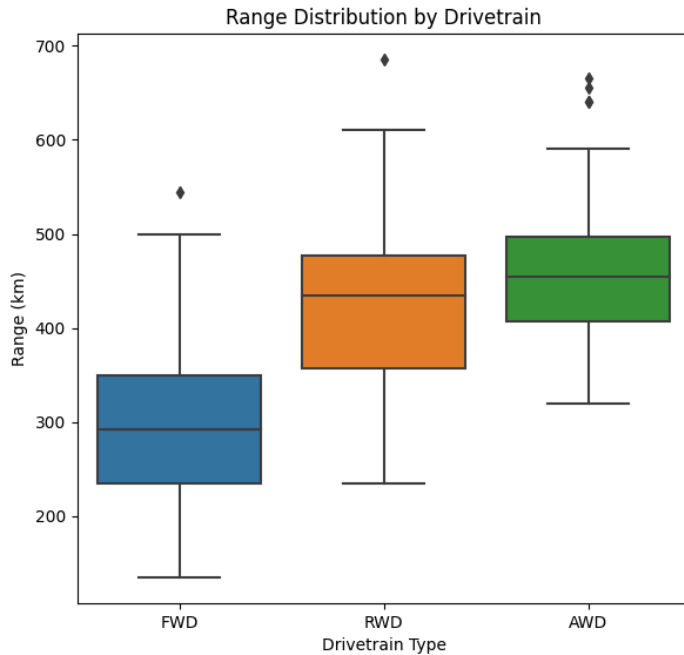
plt.tight_layout()
plt.show()

```

```

Range statistics by drivetrain:
      mean  median   std  count
drivetrain
AWD      453.77  455.0  68.33   191
FWD      298.01  292.5  78.79   156
RWD      418.17  435.0  89.60   131

```



✓ How does vehicle efficiency correlate with range?

```

df_eff = df.dropna(subset=['efficiency_wh_per_km', 'range_km'])
df_eff = df_eff[(df_eff['efficiency_wh_per_km'] < 500) & (df_eff['range_km'] < 1000)]

correlation = df_eff['efficiency_wh_per_km'].corr(df_eff['range_km'])
print(f"Correlation between efficiency and range: {correlation:.3f}")

plt.figure(figsize=(12, 5))


plt.scatter(df_eff['efficiency_wh_per_km'], df_eff['range_km'], alpha=0.6, color='brown')
plt.xlabel('Efficiency (Wh/km)')
plt.ylabel('Range (km)')
plt.title('Efficiency vs Range')

sns.regplot(data=df_eff, x='efficiency_wh_per_km', y='range_km', scatter_kws={'alpha':0.6})

plt.tight_layout()
plt.show()

```

Correlation between efficiency and range: 0.023



✓ What is the distribution of top speeds?

```
df_speed = df.dropna(subset=['top_speed_kmh'])
df_speed = df_speed[df_speed['top_speed_kmh'] < 300] # Remove outliers
```

```
print(f"Top speed statistics:")
print(f"Mean: {df_speed['top_speed_kmh'].mean():.1f} km/h")
print(f"Median: {df_speed['top_speed_kmh'].median():.1f} km/h")
print(f"Standard deviation: {df_speed['top_speed_kmh'].std():.1f} km/h")
```

```
plt.figure(figsize=(15, 5))
```

```
plt.subplot(1, 3, 1)
plt.hist(df_speed['top_speed_kmh'], bins=30, color='skyblue', alpha=0.7)
plt.xlabel('Top Speed (km/h)')
plt.ylabel('Frequency')
plt.title('Distribution of Top Speeds')
```


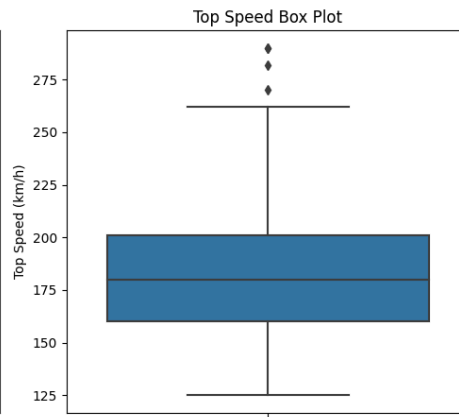
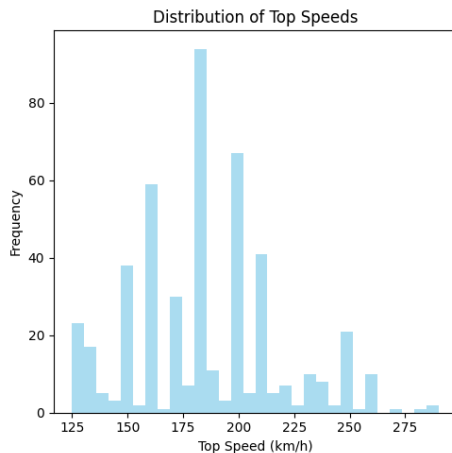
```
plt.subplot(1, 3, 2)
sns.boxplot(y=df_speed['top_speed_kmh'])
plt.ylabel('Top Speed (km/h)')
plt.title('Top Speed Box Plot')
```

```
plt.subplot(1, 3, 3)
sns.violinplot(y=df_speed['top_speed_kmh'])
plt.ylabel('Top Speed (km/h)')
plt.title('Top Speed Violin Plot')
```

```
plt.tight_layout()
plt.show()
```



Top speed statistics:
Mean: 184.9 km/h
Median: 180.0 km/h
Standard deviation: 33.3 km/h



✓ How do luxury brands compare in terms of range?

```
luxury_brands = ['Tesla', 'Mercedes', 'BMW', 'Audi', 'Porsche', 'Jaguar', 'Lucid', 'Rivian', 'Genesis']
df_luxury = df[df['brand'].isin(luxury_brands) & df['range_km'].notna()]
df_luxury = df_luxury[df_luxury['range_km'] < 1000]
```

```
print("Range statistics for luxury brands:")
luxury_range = df_luxury.groupby('brand')['range_km'].agg(['mean', 'median', 'count']).round(2)
print(luxury_range)
```

```
plt.figure(figsize=(15, 6))
```

```
plt.subplot(1, 2, 1)
sns.boxplot(data=df_luxury, x='brand', y='range_km')
plt.title('Range Distribution for Luxury Brands')
plt.xlabel('Brand')
plt.ylabel('Range (km)')
plt.xticks(rotation=45)
```

```
plt.subplot(1, 2, 2)
luxury_range_mean = df_luxury.groupby('brand')['range_km'].mean().sort_values(ascending=False)
```

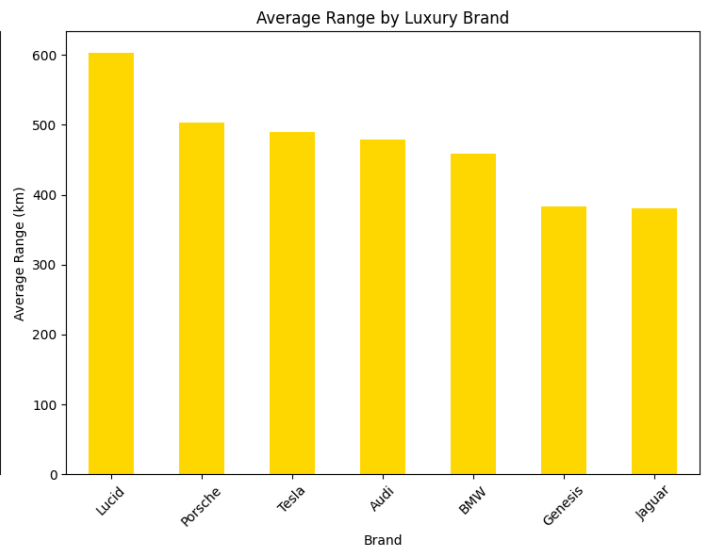
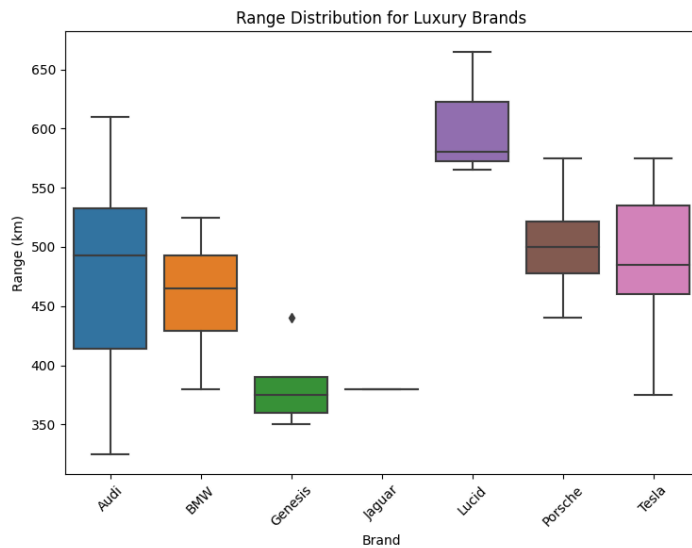


```
luxury_range_mean.plot(kind='bar', color='gold')
plt.title('Average Range by Luxury Brand')
plt.xlabel('Brand')
plt.ylabel('Average Range (km)')
plt.xticks(rotation=45)
```

```
plt.tight_layout()
plt.show()
```

↔ Range statistics for luxury brands:

	mean	median	count
brand			
Audi	478.39	492.5	28
BMW	458.25	465.0	20
Genesis	383.00	375.0	5
Jaguar	380.00	380.0	1
Lucid	603.33	580.0	3
Porsche	502.88	500.0	26
Tesla	490.00	485.0	11



✓ Relationship between vehicle dimensions and acceleration

```
df_dim = df.dropna(subset=['length_mm', 'acceleration_0_100_s'])
df_dim = df_dim[df_dim['acceleration_0_100_s'] < 15]

correlation = df_dim['length_mm'].corr(df_dim['acceleration_0_100_s'])
print(f"Correlation between length and acceleration: {correlation:.3f}")

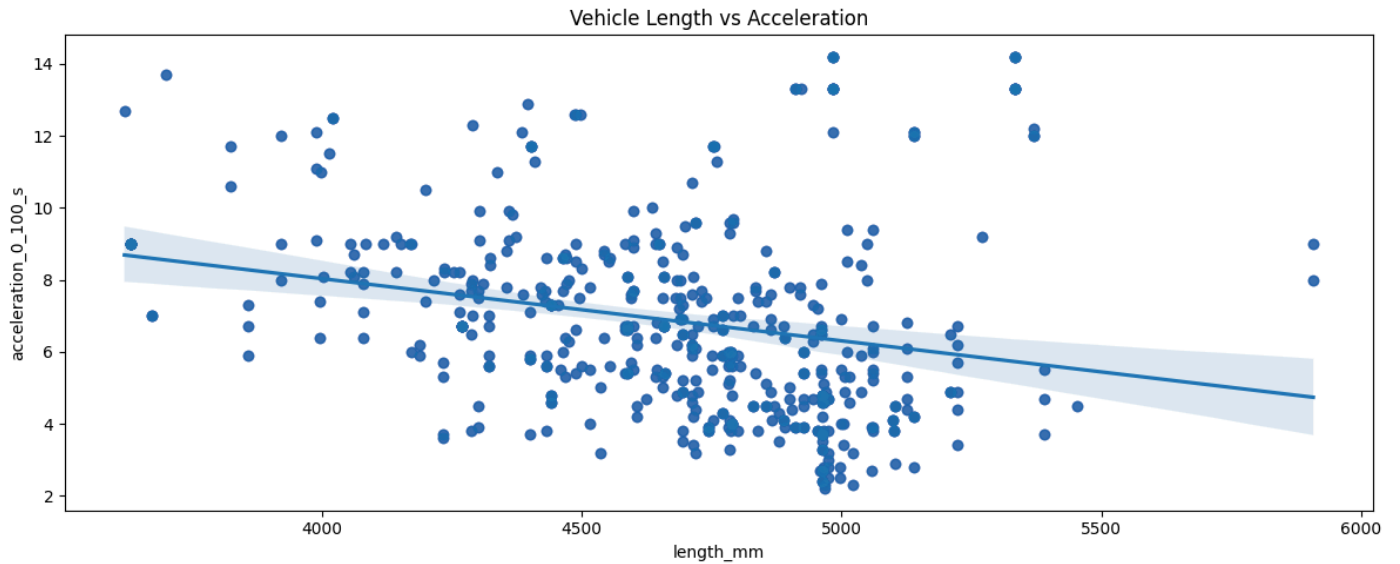
plt.figure(figsize=(12, 5))

plt.scatter(df_dim['length_mm'], df_dim['acceleration_0_100_s'], alpha=0.6, color='navy')
plt.xlabel('Vehicle Length (mm)')
plt.ylabel('0-100 km/h (seconds)')
plt.title('Vehicle Length vs Acceleration')

sns.regplot(data=df_dim, x='length_mm', y='acceleration_0_100_s', scatter_kws={'alpha':0.6})

plt.tight_layout()
plt.show()
```

Correlation between length and acceleration: -0.237



Seating capacity distribution

```
df_seats = df.dropna(subset=['seats'])
df_seats = df_seats[df_seats['seats'] <= 10] # Remove outliers

seat_counts = df_seats['seats'].value_counts().sort_index()
print("Seating capacity distribution:")
print(seat_counts)

plt.figure(figsize=(12, 6))

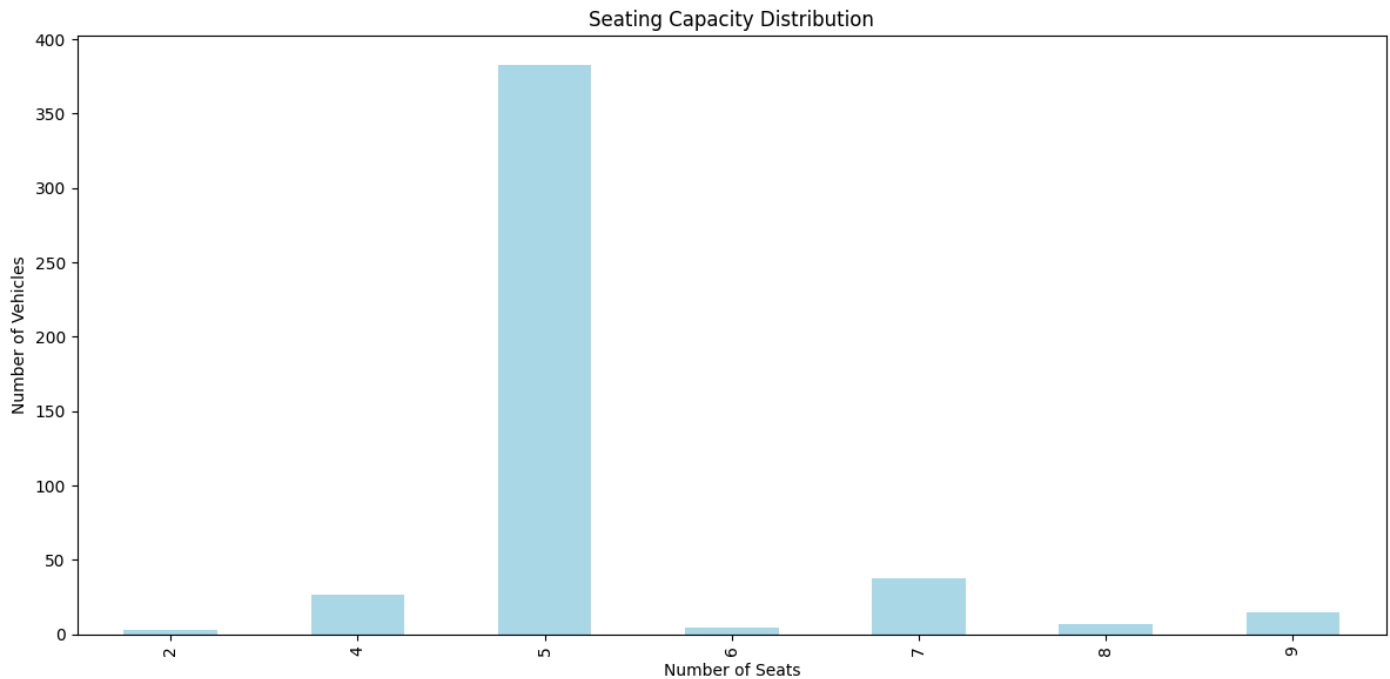
seat_counts.plot(kind='bar', color='lightblue')
plt.title('Seating Capacity Distribution')
plt.xlabel('Number of Seats')
plt.ylabel('Number of Vehicles')

plt.tight_layout()
plt.show()
```

```

Seating capacity distribution:
seats
2      3
4     27
5    383
6      5
7     38
8      7
9     15
Name: count, dtype: int64

```



Towing capacity by segment

```

df_tow = df.dropna(subset=['towing_capacity_kg', 'segment'])
df_tow = df_tow[df_tow['towing_capacity_kg'] > 0] # Only vehicles with towing capacity

print("Towing capacity statistics by segment:")
towing_stats = df_tow.groupby('segment')['towing_capacity_kg'].agg(['mean', 'median', 'count']).round(2)
print(towing_stats)

plt.figure(figsize=(15, 6))

plt.subplot(1, 2, 1)
sns.boxplot(data=df_tow, x='segment', y='towing_capacity_kg')
plt.title('Towing Capacity by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Towing Capacity (kg)')
plt.xticks(rotation=45)

plt.subplot(1, 2, 2)
towing_mean = df_tow.groupby('segment')['towing_capacity_kg'].mean().sort_values(ascending=False)
towing_mean.plot(kind='bar', color='red')
plt.title('Average Towing Capacity by Segment')
plt.xlabel('Vehicle Segment')
plt.ylabel('Average Towing Capacity (kg)')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

```